



BERT, Compression and Applications

Dongkuan (DK) Xu, Ph.D. Candidate

Pennsylvania State University, Advisor: Xiang Zhang

Web: www.personal.psu.edu/dux19/


04-08-2021

Agenda

- BERT Architecture
- Model Compression
- Applications

Background

- BERT (Bidirectional Encoder Representations from Transformers)
 - Published by Google AI Language [1]
 - Achieved state-of-the-art results in various NLP/CV tasks
- Key Innovation
 - Bidirectional training for language modelling
 - Previous efforts looked at a text sequence from left-to-right, right-to-left or combined way



Background

- Performance on GLUE (General Language Understanding Evaluation) Benchmark
 - The most popular collection for training, evaluating and analyzing NLP systems [1]
 - Constructed by NYU, UW and DeepMind

gluebenchmark.com

GLUE

SuperGLUE

Paper

</>

Code

Tasks

Leaderboard

FAQ

Diagnostics

Submit

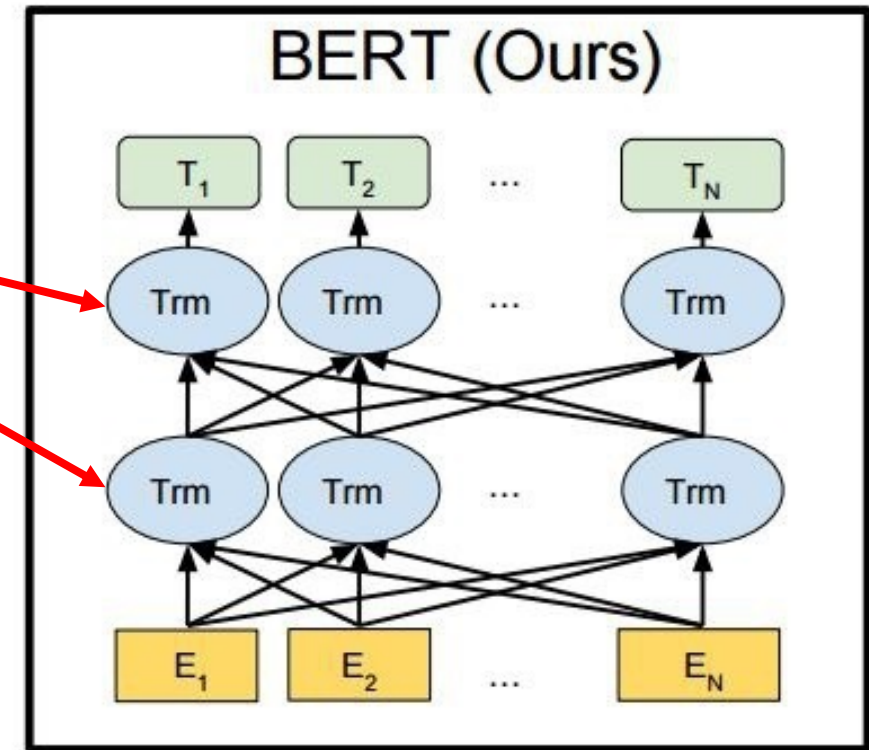
Login

Rank	Name	Model	URL	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI	RTE	WNLI	AX	
+	1	PING-AN Omni-Sinitic	ALBERT + DAAF + NAS		90.6	73.5	97.2	94.0/92.0	93.0/92.4	76.1/91.0	91.6	91.3	97.5	91.7	94.5	51.2
	2	ERNIE Team - Baidu	ERNIE	↗	90.4	74.4	97.5	93.5/91.4	93.0/92.6	75.2/90.9	91.4	91.0	96.6	90.9	94.5	51.7
+	3	Alibaba DAMO NLP	StructBERT	↗	90.3	75.3	97.1	93.9/91.9	93.0/92.5	74.8/91.0	90.9	90.7	96.4	90.2	94.5	49.1
	4	T5 Team - Google	T5	↗	90.3	71.6	97.5	92.8/90.4	93.1/92.8	75.1/90.6	92.2	91.9	96.9	92.8	94.5	53.1
	5	Microsoft D365 AI & MSR AI & GATECH	MT-DNN-SMART	↗	89.9	69.5	97.5	93.7/91.6	92.9/92.5	73.9/90.2	91.0	90.8	99.2	89.7	94.5	50.2
+	6	ELECTRA Team	ELECTRA-Large + Standard Tricks	↗	89.4	71.7	97.1	93.1/90.7	92.9/92.5	75.6/90.8	91.3	90.8	95.8	89.8	91.8	50.7
+	7	Huawei Noah's Ark Lab	NEZHA-Large		88.7	67.4	97.2	93.2/91.0	92.2/91.6	74.1/90.2	90.8	90.2	95.7	88.5	93.2	45.0
+	8	Microsoft D365 AI & UMD	FreeLB-RoBERTa (ensemble)	↗	88.4	68.0	96.8	93.1/90.8	92.3/92.1	74.8/90.3	91.1	90.7	95.6	88.7	89.0	50.1
	9	Junjie Yang	HIRE-RoBERTa	↗	88.3	68.6	97.1	93.0/90.7	92.4/92.0	74.3/90.2	90.7	90.4	95.5	87.9	89.0	49.3
	10	Facebook AI	RoBERTa	↗	88.1	67.8	96.7	92.3/89.8	92.2/91.9	74.3/90.2	90.8	90.2	95.4	88.2	89.0	48.7

Support BERT
-> Support All

Overall

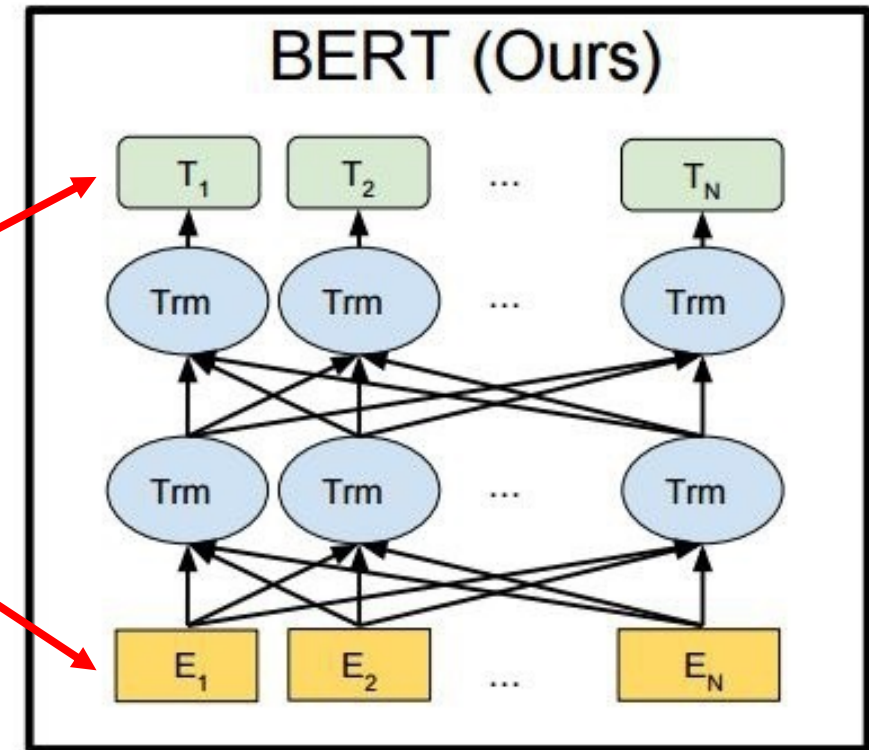
- General architecture
 - Multiple Transformer encoders
 - Input: Embeddings of words
 - Output: Hidden representations of words
- Downstream task
 - e.g., sentence classification
- How BERT works
 - Pre-training
 - The model is trained on unlabeled data over different pre-training tasks.
 - Fine-tuning
 - The model is first initialized with the pre-trained parameters, and all the parameters are fine-tuned using labeled data from the downstream tasks.



Architecture of BERT

Overall

- General architecture
 - Multiple Transformer encoders
 - Input: Embeddings of words
 - Output: Hidden representations of words
- Downstream task
 - e.g., sentence classification
- How BERT works
 - Pre-training
 - The model is trained on unlabeled data over different pre-training tasks.
 - Fine-tuning
 - The model is first initialized with the pre-trained parameters, and all the parameters are fine-tuned using labeled data from the downstream tasks.

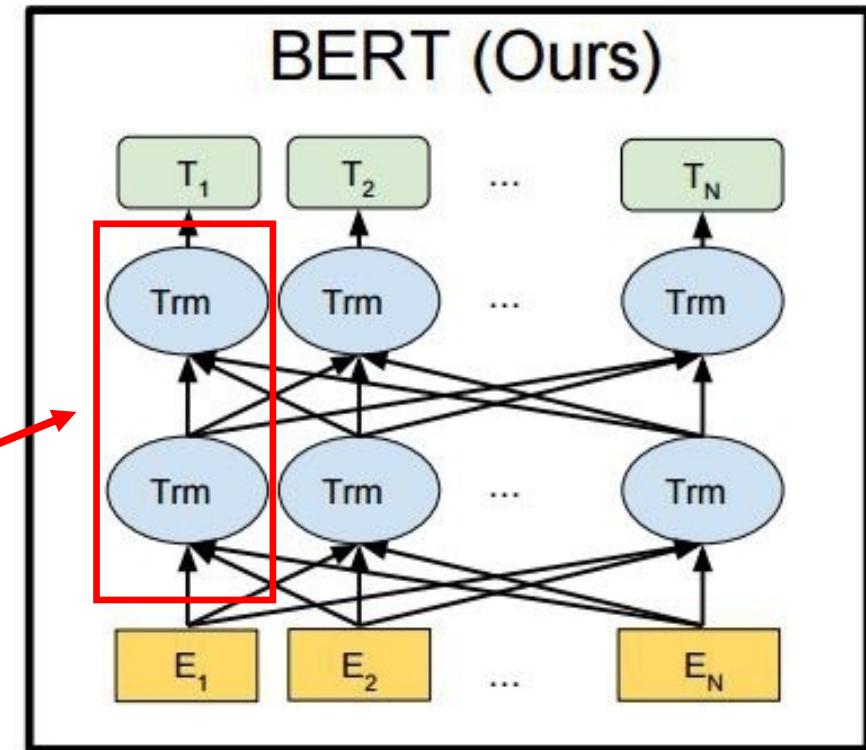


Architecture of BERT

Overall

- General architecture
 - Multiple Transformer encoders
 - Input: Embeddings of words
 - Output: Hidden representations of words
- Downstream task
 - e.g., sentence classification
- How BERT works
 - Pre-training
 - The model is trained on unlabeled data over different pre-training tasks.
 - Fine-tuning
 - The model is first initialized with the pre-trained parameters, and all the parameters are fine-tuned using labeled data from the downstream tasks.

- 1. Only one series of encoders
- 2. Shared by all word embeddings



Architecture of BERT

Overall

- General architecture
 - Multiple Transformer encoders
 - Input: Embeddings of words
 - Output: Hidden representations of words
- Downstream task
 - e.g., sentence classification
- How BERT works
 - Pre-training
 - The model is trained on unlabeled data over different pre-training tasks.
 - Fine-tuning
 - The model is first initialized with the pre-trained parameters, and all the parameters are fine-tuned using labeled data from the downstream tasks.

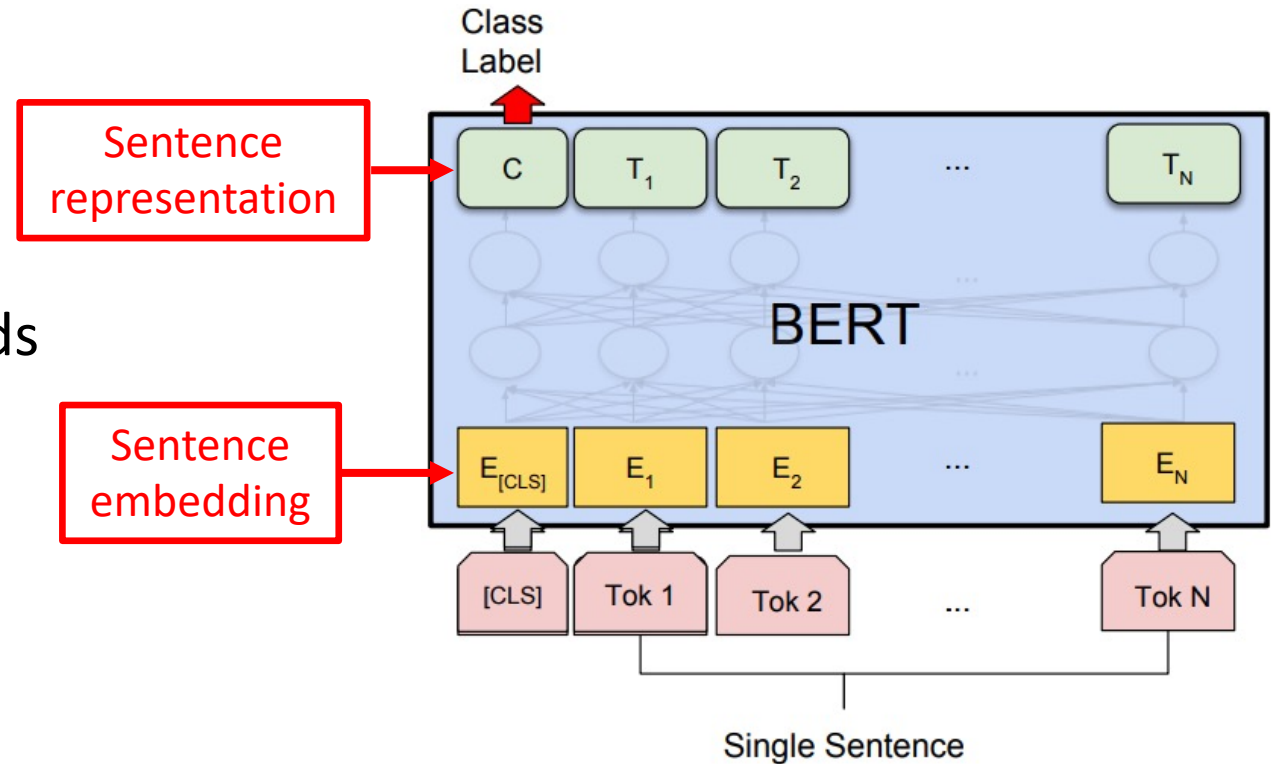


Illustration of BERT on classification

Overall

- General architecture
 - Multiple Transformer encoders
 - Input: Embeddings of words
 - Output: Hidden representations of words

- Downstream task
 - e.g., sentence classification

- How BERT works

- Pre-training

- The model is trained on unlabeled data over different pre-training tasks.

- Fine-tuning

- The model is first initialized with the pre-trained parameters, and all the parameters are fine-tuned using labeled data from the downstream tasks.

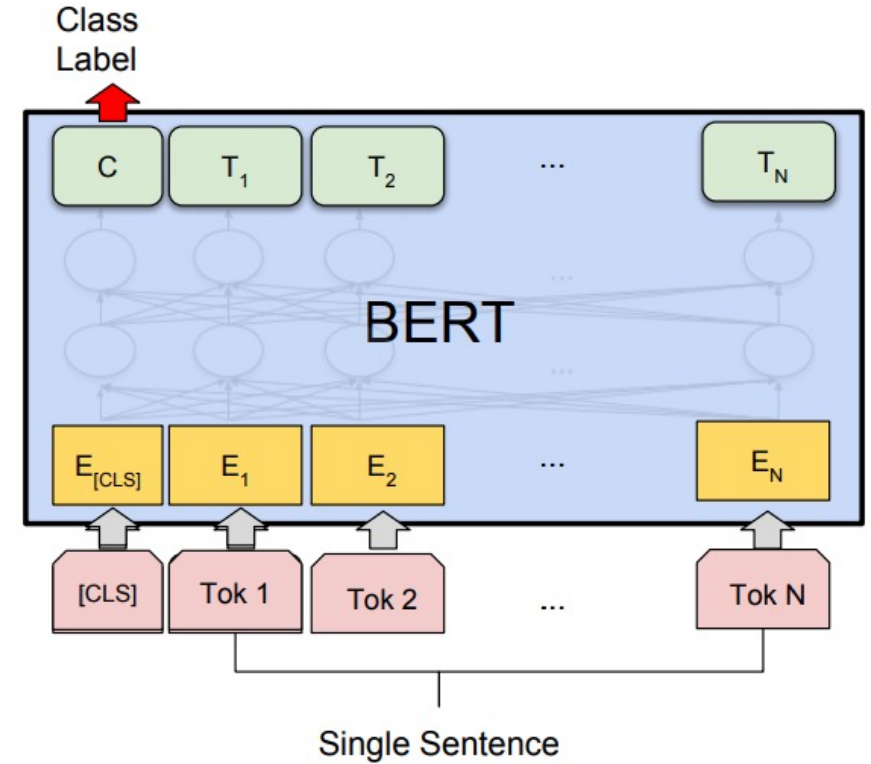

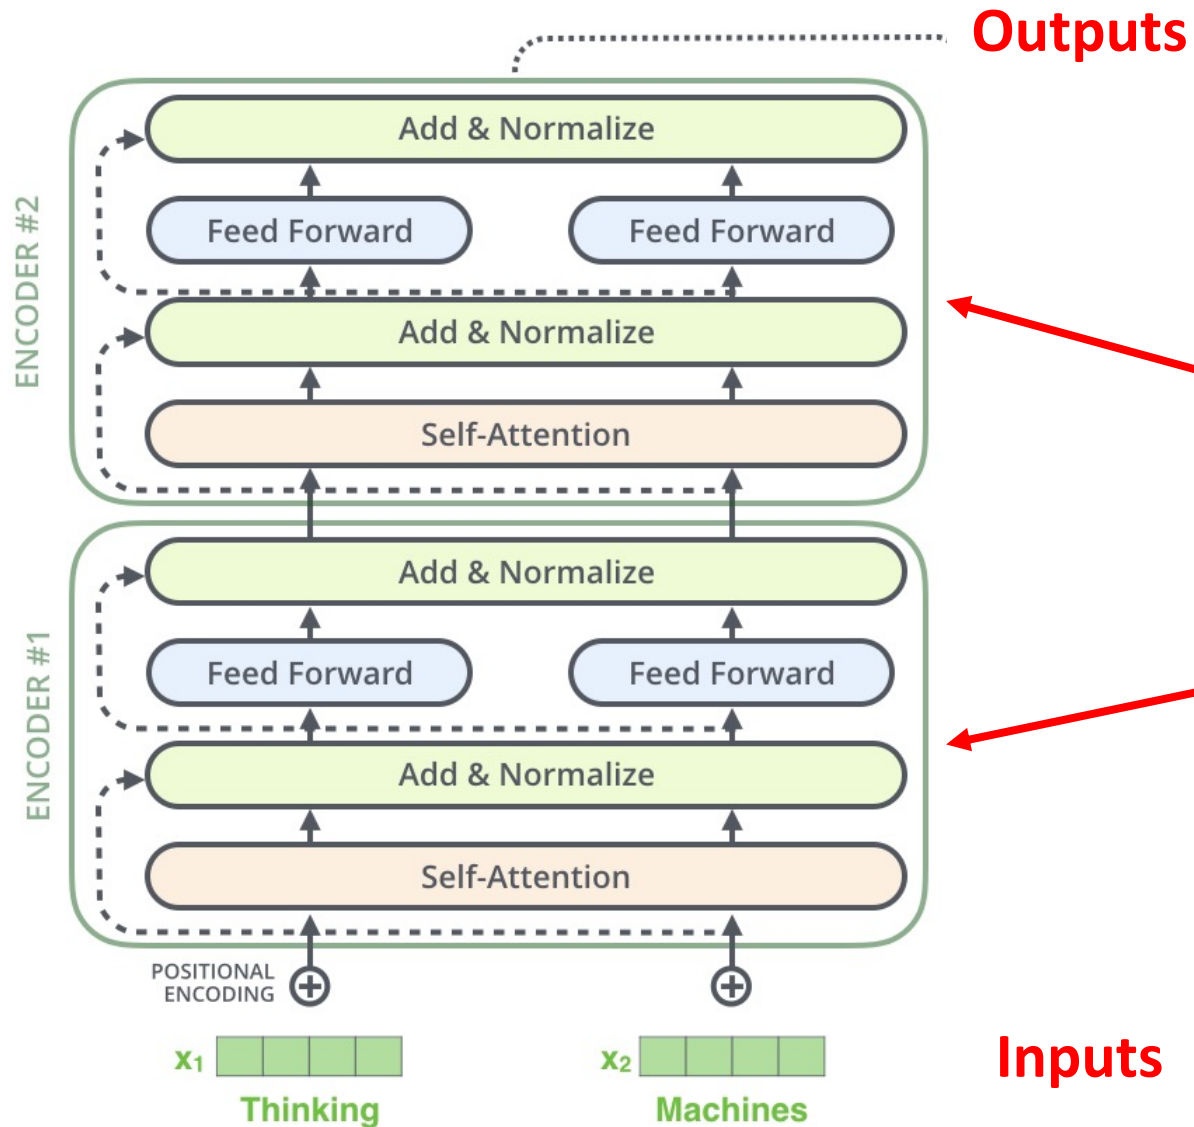


Illustration of BERT on classification

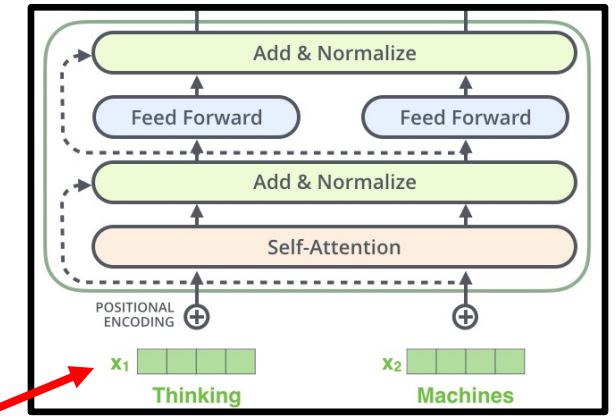
Transformer Encoder



Architecture of Encoder (Two Encoders Here)

Embedding Look_Up

- Embedding is the element-wise sum of three embeddings
- Goal
 - To give the model a sense of the order of the words



Architecture of Transformer Encoder

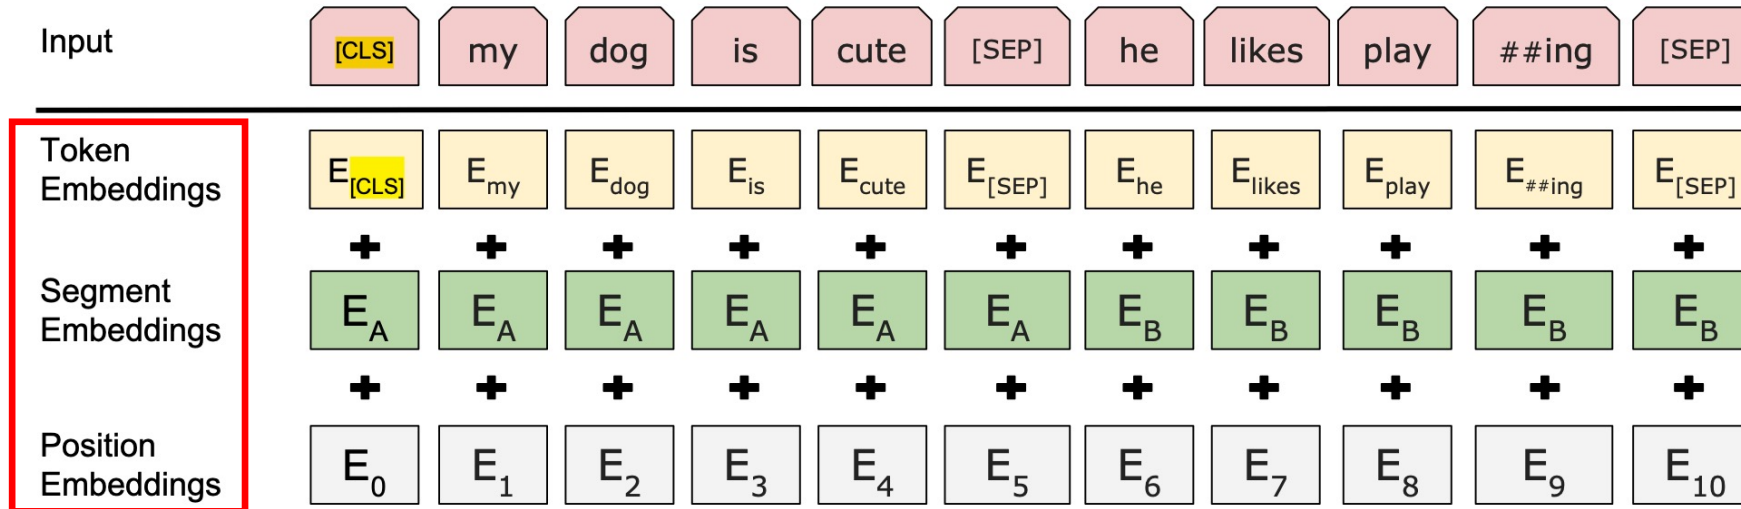


Illustration of BERT Input Representation

Embedding Look_Up

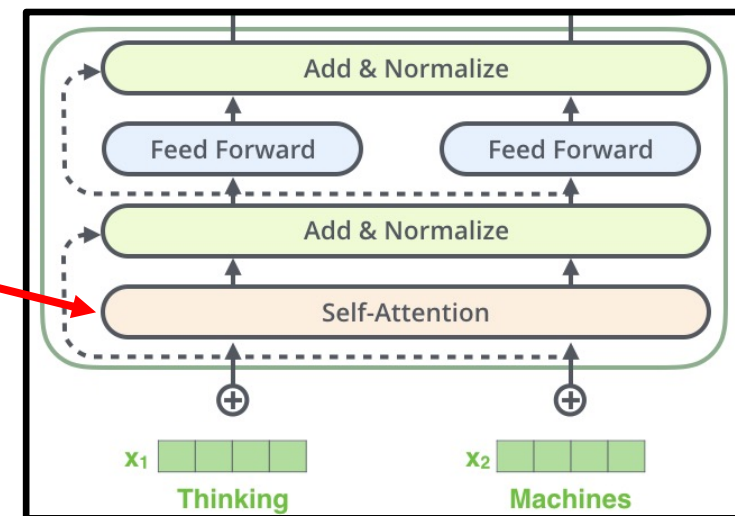
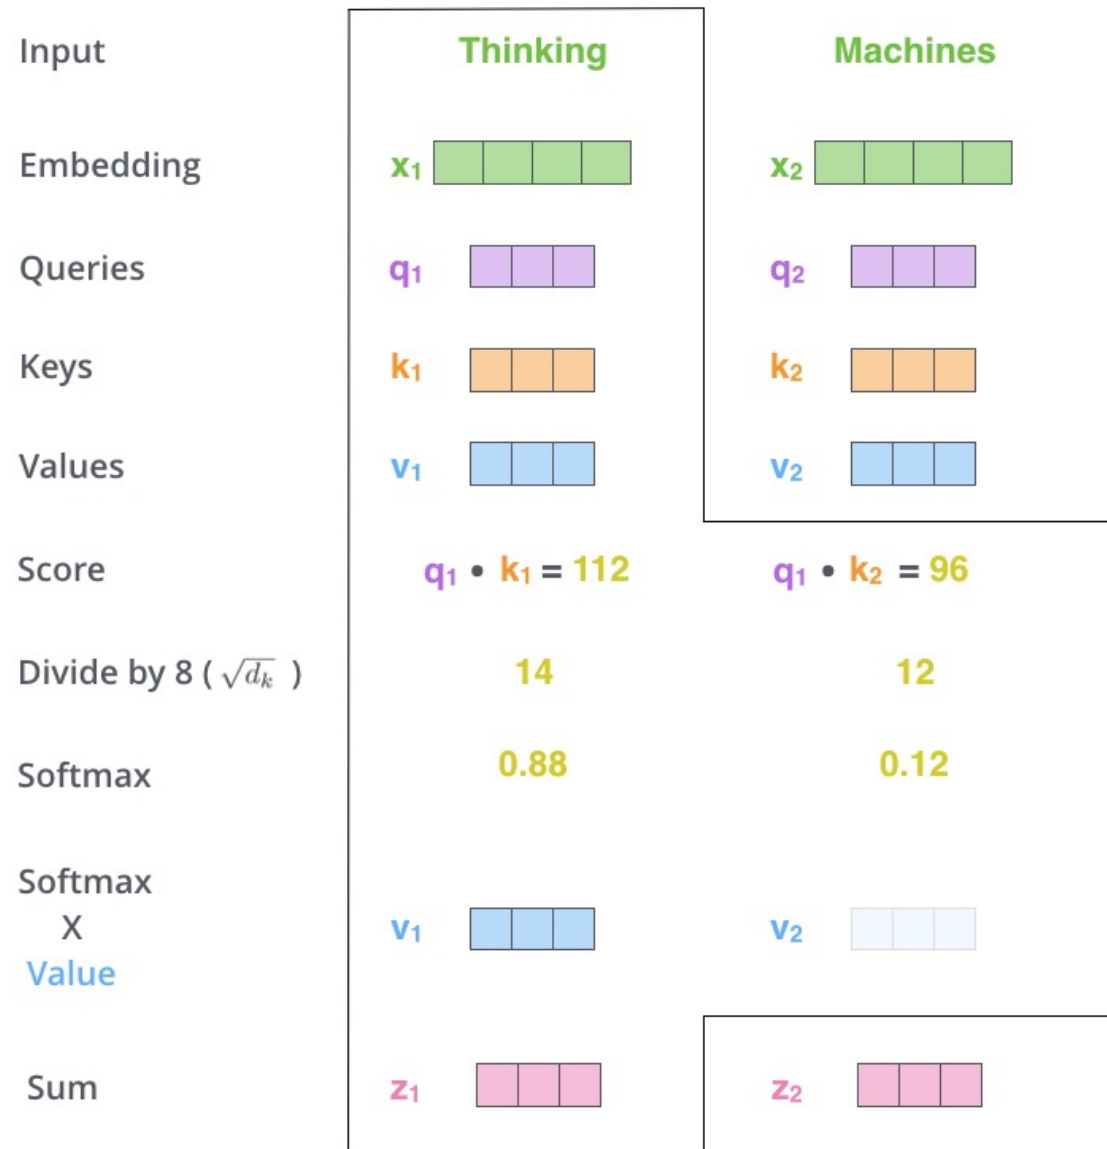
- Huge embedding table lookup

```
3     class BertEmbeddings(nn.Module):  
4         """Construct the embeddings from word, position and token_type embeddings.  
5         """  
6         def __init__(self, config):  
7             super(BertEmbeddings, self).__init__()  
8             self.word_embeddings = nn.Embedding(config.vocab_size, config.hidden_size)  
9             self.position_embeddings = nn.Embedding(config.max_position_embeddings, config.hidden_size)  
10            self.token_type_embeddings = nn.Embedding(config.type_vocab_size, config.hidden_size)
```

- `nn.Embedding()`
 - A simple lookup table that stores embeddings
 - All the elements are parameters
 - Input is a list of indices, and output is the word embeddings
- Implementation links
 - PyTorch: https://pytorch.org/docs/stable/_modules/torch/nn/modules/sparse.html#Embedding
 - TensorFlow: https://github.com/tensorflow/tensorflow/blob/v2.2.0/tensorflow/python/ops/embedding_ops.py#L329-L373

Self-Attention

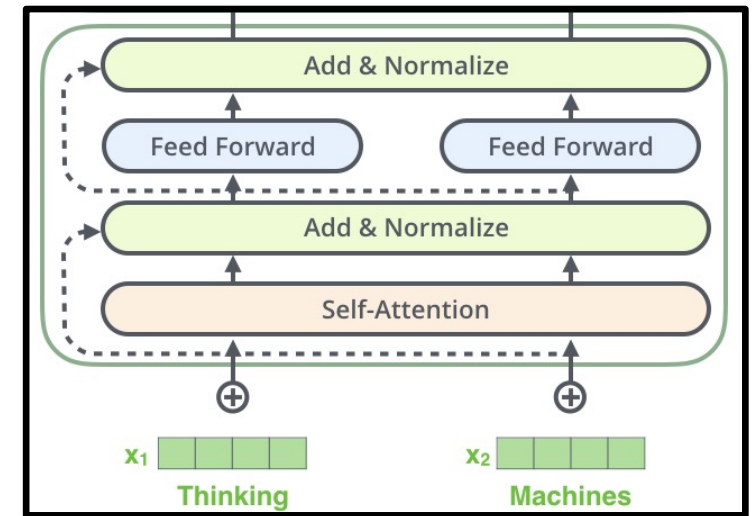
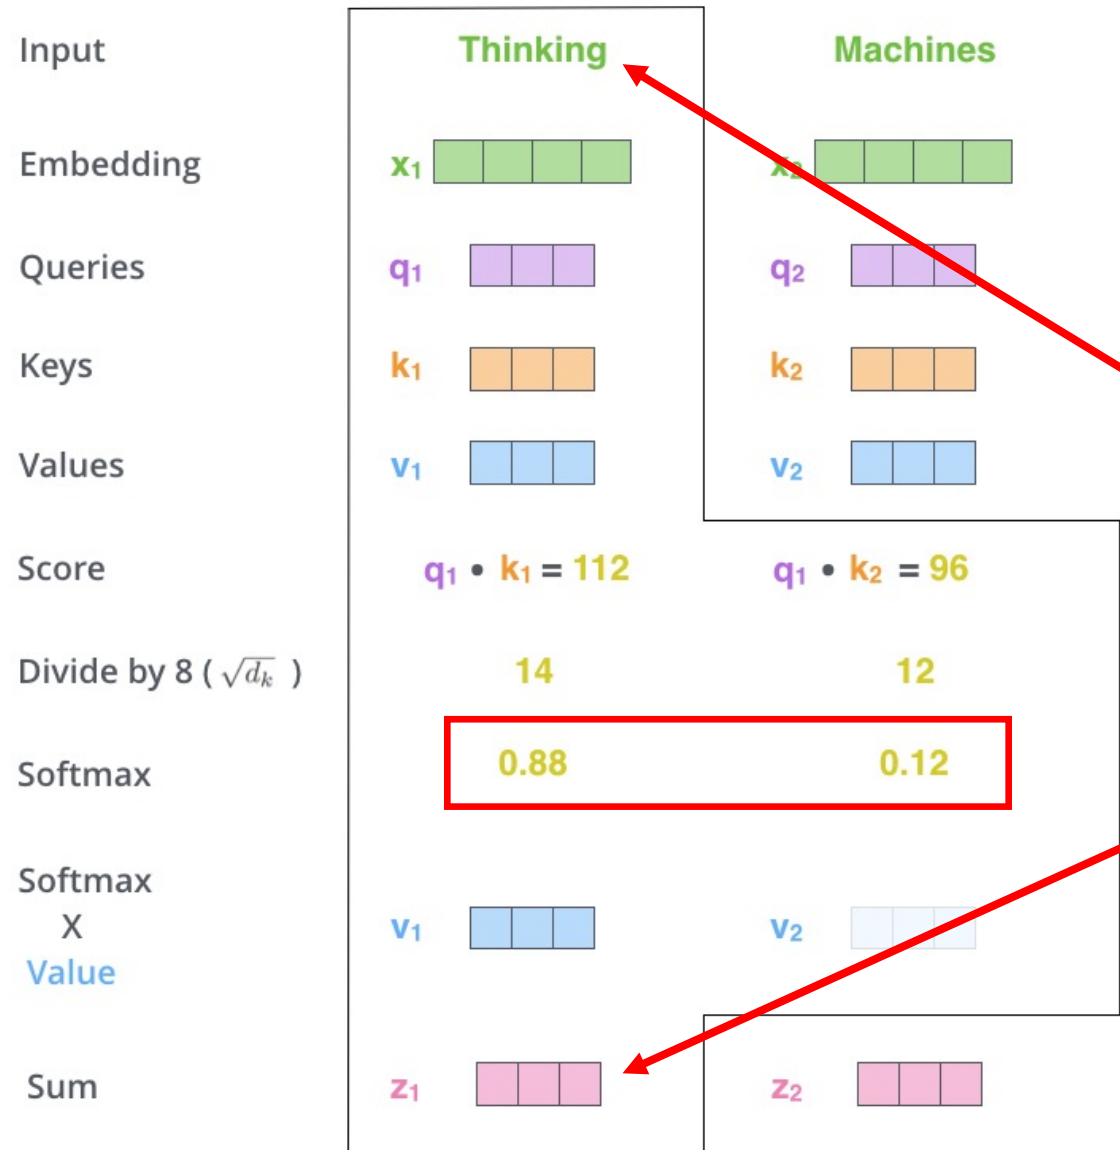
- Take two words as an example



Architecture of Transformer Encoder

Self-Attention

- Take two words as an example




Architecture of Transformer Encoder

Input: Word embedding vectors
Output: New vector representations

Self-Attention in Matrix Calculation

- First step to calculate Q, K, V

- Rows: **words of a sentence**
- Columns: **hidden_dims**



- Second step to calculate the output Z

- Rows: **words of a sentence**
- Columns: **hidden_dims**

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}}\right) \times V$$

= Z

The diagram illustrates the second step of self-attention: calculating the output matrix Z. The query matrix Q (2x3 purple grid) is multiplied by the transpose of the key matrix K^T (3x2 orange grid). The result is divided by the square root of the key dimension $\sqrt{d_k}$. The softmax function is applied to the result. The final output matrix Z is a 2x3 grid of pink squares, highlighted by a red box.

Multi-Head Self-Attention

1) This is our input sentence*

2) We embed each word*



3) Split into 8 heads. We multiply X or R with weight matrices

4) Calculate attention using the resulting $Q/K/V$ matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

Thinking
Machines




X



One attention head

* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



R



...

...


...



W^O



Z



Layer Normalization [1]

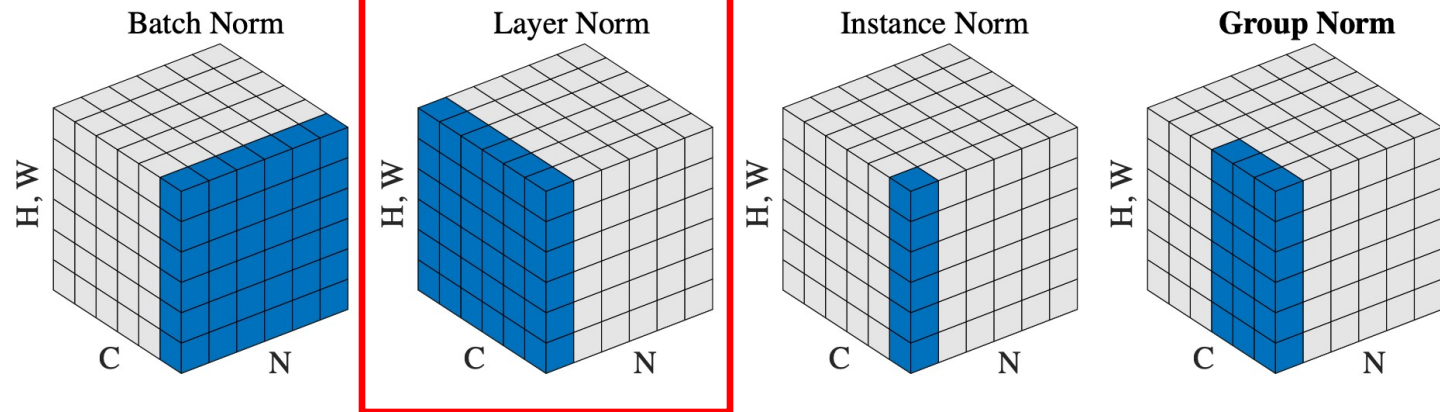
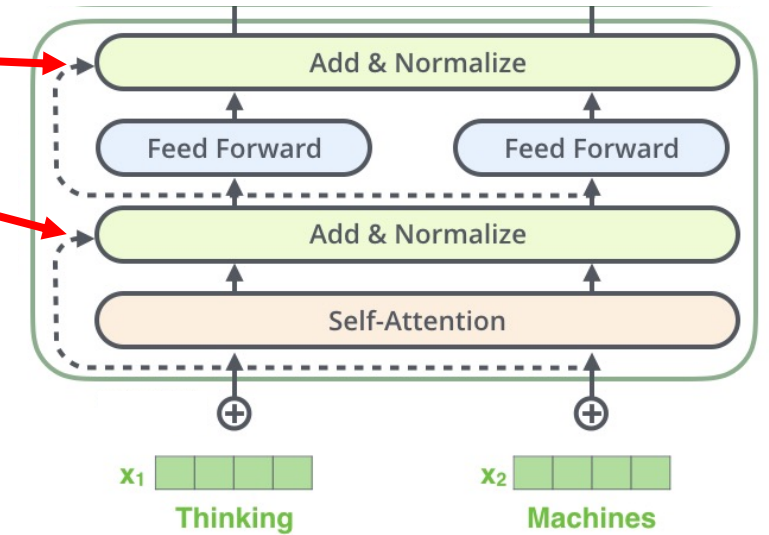
- Motivations

- 1. Dynamic length
- 2. Different meaning of the same position

- Formula

$$y = \frac{x - \mathbf{E}[x]}{\sqrt{\mathbf{Var}[x] + \epsilon}} * \gamma + \beta$$

- Comparison between four normalizations [2]



[1] Ba, J. L., Kiros, J. R., & Hinton, G. E. (2016). Layer normalization. *arXiv preprint arXiv:1607.06450*.

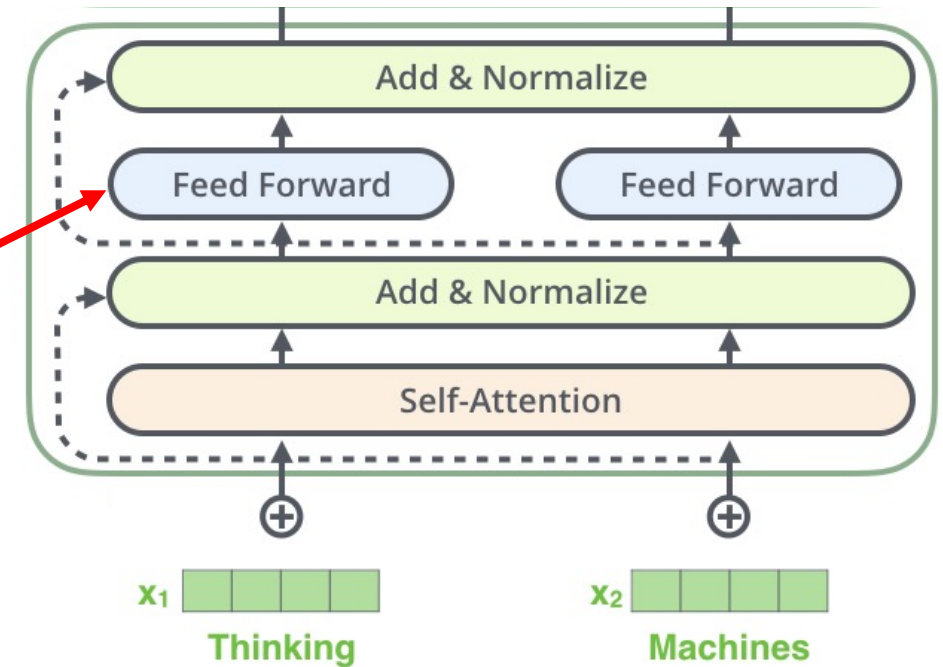
[2] Wu, Y., & He, K. (2018). Group normalization. In *Proceedings of the European Conference on Computer Vision (ECCV)* (pp. 3-19).

Feed-Forward Networks

- Architecture
 - Two linear transformations with a ReLU in between
 - Dimension of input and output = 512
 - Dimension of inner-layer = 512×4

- Formula

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$



Other Transformer Based NLP Models

- **XLNet (CMU + Google AI)**
 - Yang, Z., Dai, Z., Yang, Y., Carbonell, J., Salakhutdinov, R. R., & Le, Q. V. (2019). Xlnet: Generalized autoregressive pretraining for language understanding. In *Advances in neural information processing systems* (pp. 5754-5764).
- **ALBERT (Google Language)**
 - Lan, Z., Chen, M., Goodman, S., Gimpel, K., Sharma, P., & Soricut, R. (2019). Albert: A lite bert for self-supervised learning of language representations. arXiv preprint arXiv:1909.11942.
- **RoBERTa (Facebook AI)**
 - Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., ... & Stoyanov, V. (2019). Roberta: A robustly optimized bert pretraining approach. arXiv preprint arXiv:1907.11692.
- **Transformer-XL (CMU + Google Brain)**
 - Dai, Z., Yang, Z., Yang, Y., Carbonell, J., Le, Q. V., & Salakhutdinov, R. (2019). Transformer-xl: Attentive language models beyond a fixed-length context. arXiv preprint arXiv:1901.02860.
- **ERNIE (Baidu)**
 - Sun, Y., Wang, S., Li, Y., Feng, S., Chen, X., Zhang, H., ... & Wu, H. (2019). Ernie: Enhanced representation through knowledge integration. *arXiv preprint arXiv:1904.09223*.
- **GPT-2 (OpenAI)**
 - Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language models are unsupervised multitask learners. OpenAI Blog, 1(8), 9.

Support BERT -> Support All

Agenda

- BERT Architecture
- Model Compression
- Applications

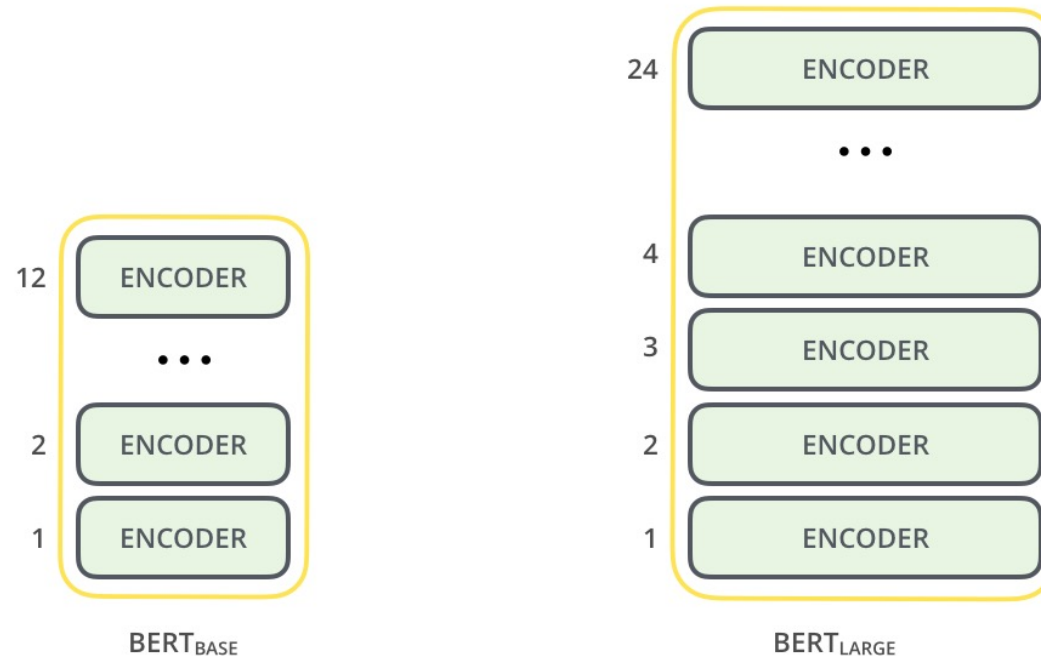
BERT_{Base} v.s. BERT_{Large}

- BERT_{Base}

- #para = 110M, #encoders = 12, #dim = 768, #head = 12,
- #FLOPs = 123M * sentence_length * #Batch

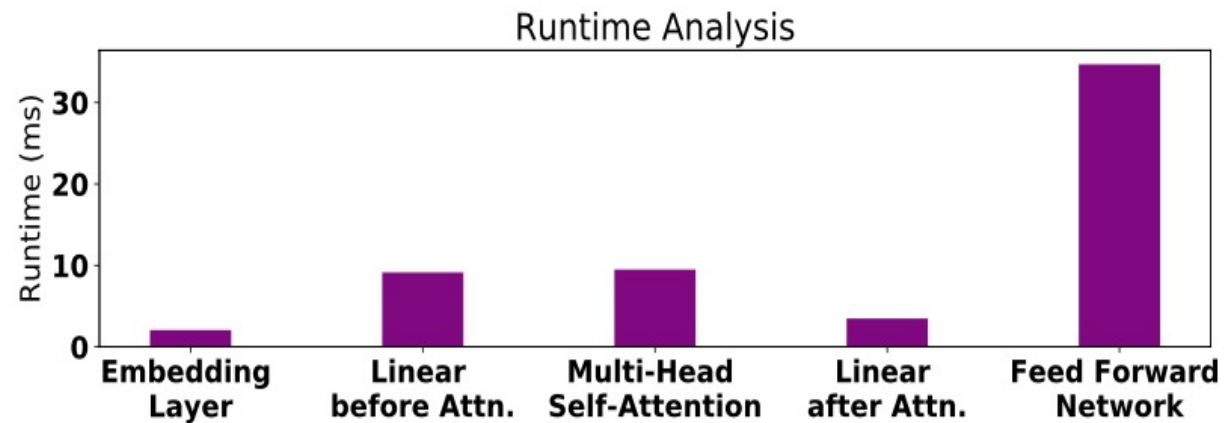

- BERT_{Large}

- #para = 340M, #encoders = 24, #dim = 1024, #head = 16
- #FLOPs = 1857M * sentence_length * #Batch



Comparison between BERT_{base} and BERT_{Large}

Parameter / FLOPs Distributions



Computation Analysis [1]

Compression Methods on BERT

- Data Quantization
 - Embedding layer is more sensitive to quantization than other layers
 - More bits to maintain its accuracy
- Pruning
 - **Sparse pruning** (trending and with promising future):
 - Our NAACL'21 work [1], **SparseBERT**, achieved SOTA (**compression ratio = x20, only 1.4% performance drop**)
 - Structured pruning: #encoders, #att_heads, #hidden_dims
- Knowledge Distillation
 - Distillation on output logits, encoder outputs, attention maps
- Architecture-Invariant Compression
 - Parameter sharing, weight matrix decomposition

All The Ways You Can Compress BERT [1]

- Literatures

Paper	Prune	Factor	Distill	W. Sharing	Quant.	Pre-train	Downstream
Compressing BERT: Studying the Effects of Weight Pruning on Transfer Learning	☑					☑	☑
Are Sixteen Heads Really Better than One?	☑						☑
Pruning a BERT-based Question Answering Model	☑						☑

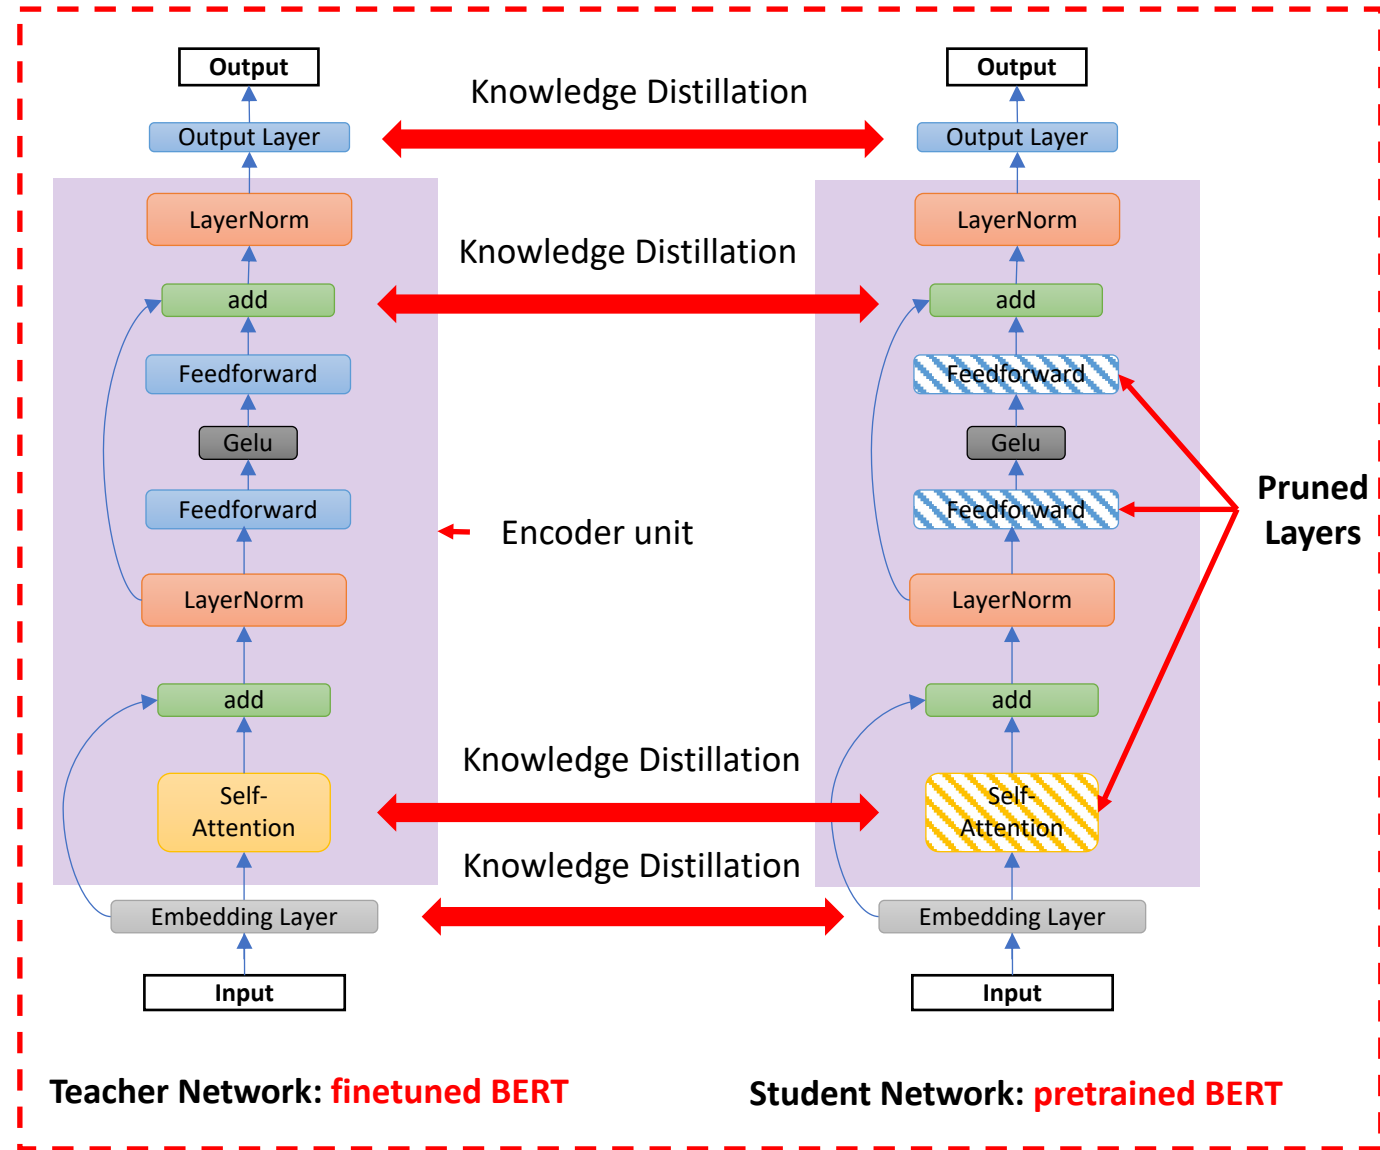
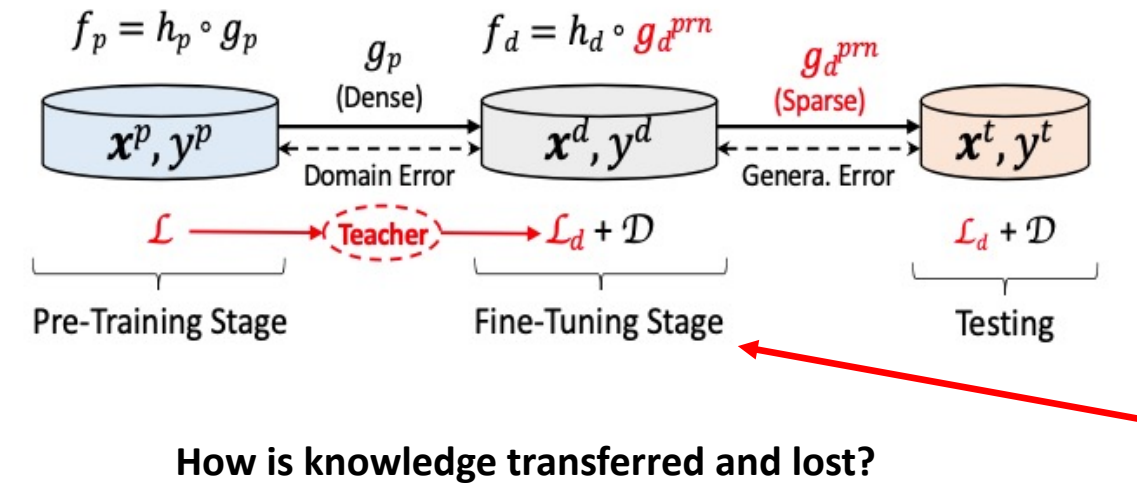
- Experimental Results

Paper	Reduction	Of	Speed-up	Accuracy?	Comments
Compressing BERT: Studying the Effects of Weight Pruning on Transfer Learning	30%	params	?	Same	Some interesting ablation experiments and fine-tuning analysis
Are Sixteen Heads Really Better than One?	50-60%	attn heads	1.2x	Same	
Pruning a BERT-based Question Answering Model	50%	attn Heads + FF	2x	-1.5 F1	

[1] <http://mitchgordon.me/machine/learning/2019/11/18/all-the-ways-to-compress-BERT.html>

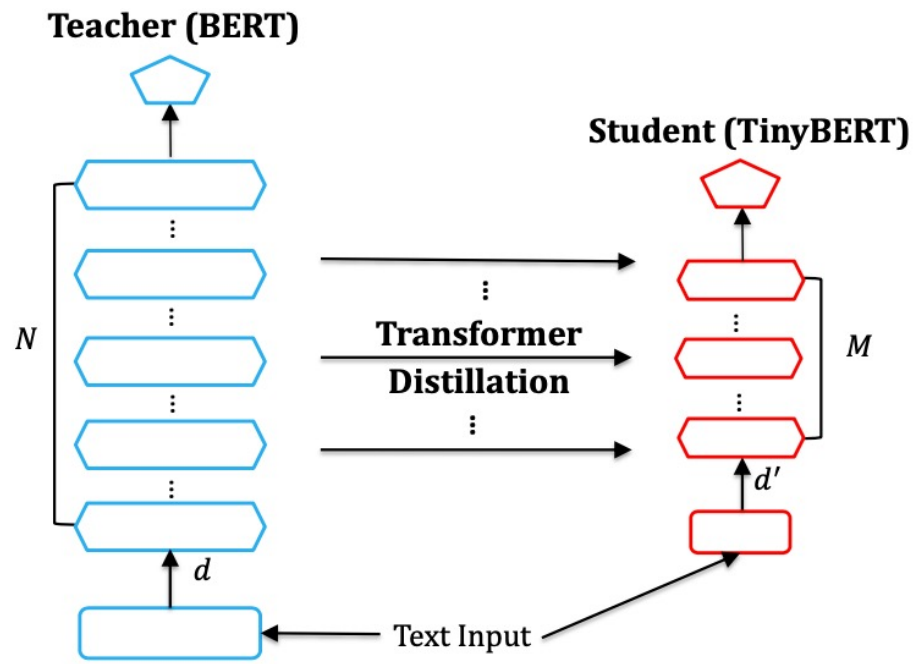
SparseBERT: Knowledge-Aware Sparse Pruning [1]

- Motivation: gap of sparse pruning in NLP
- Analysis about knowledge lost



TinyBERT: Distilling BERT for Natural Language Understanding [1]

- Knowledge distillation of the Transformer-based models
- A two-stage learning framework
- Results
 - > 96% the performance of teachers
 - 7.5x smaller and 9.4x faster on inference



Overview of TinyBERT distillation

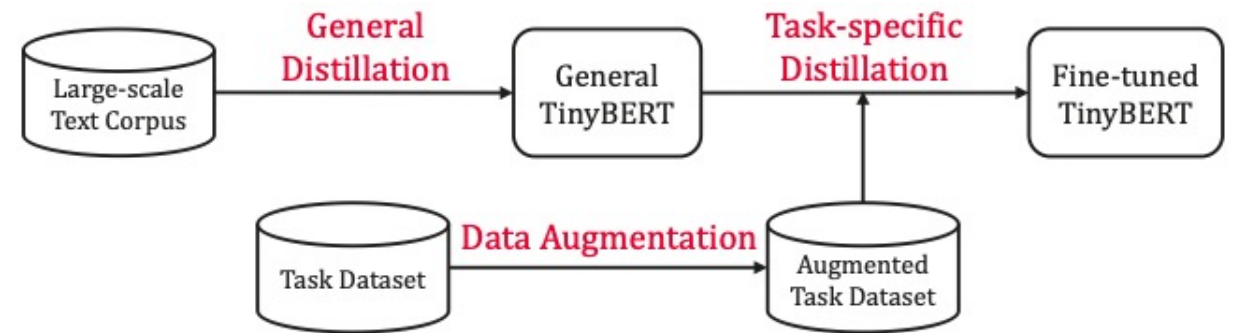






















Illustration of TinyBERT learning

[1] Huawei Noah's Ark Lab. Tinybert: Distilling bert for natural language understanding. *EMNLP 2020*.

Agenda

- BERT Architecture
- Model Compression
- Applications

Novel CV Applications using Transformers [1]

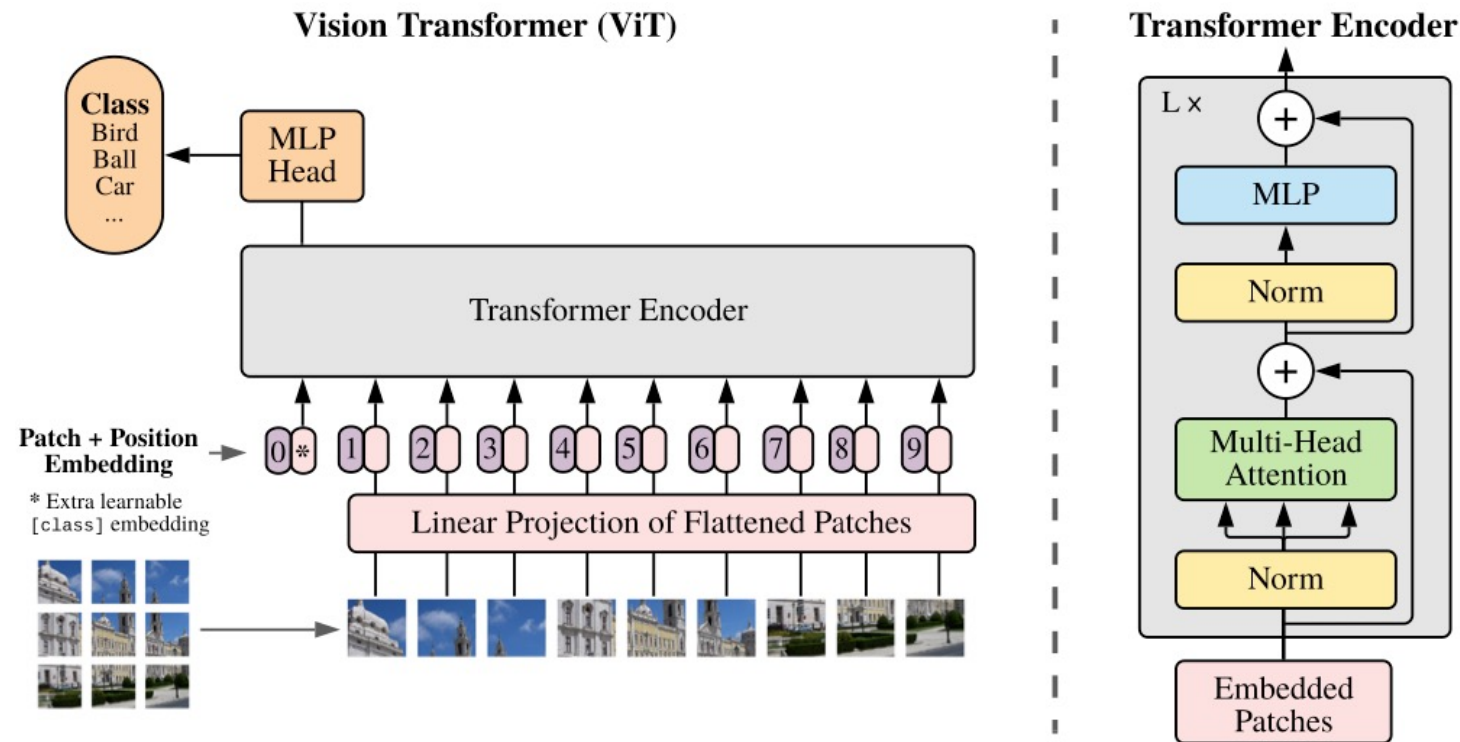
-  Transformer for Image Synthesis -  Esser et al. (2020)
-  Transformer for Multi-Object Tracking -  Sun et al. (2020)
-  Transformer for Music Generation -  Hsiao et al. (2021)
-  Transformer for Dance Generation with Music -  Huang et al. (2021)
-  Transformer for 3D Object Detection -  Bhattacharyya et al. (2021)
-  Transformer for Point-Cloud Processing -  Guo et al. (2020)
-  Transformer for Time-Series Forecasting -  Lim et al. (2020)
-  Transformer for Vision-Language Modeling -  Zhang et al. (2021)
-  Transformer for Lane Shape Prediction -  Liu et al. (2020)
-  Transformer for End-to-End Object Detection -  Zhu et al. (2021)

NLP Applications using Transformers

- ☐ Question Answer and Reading Comprehension
- ☐ Web Search and Information Retrieval
- ☐ Dialog System or Chatbot
- ☐ Text Summarization
- ☐ Data Augmentation in NLP areas
- ☐ Text/Word Classification
- ☐ Sequence Labelling
- ☐ Others

Vision Transformer (ViT) [1]


- Open the door of applying Transformer architecture to CV domains
- Split an image into fixed-size patches
- Transformer encoder



Model overview

Vision Transformer (ViT) [1]

- Vision Transformer >> SOTA convolutional networks
 - Need to be pre-trained on large data and transferred to smaller ones
- Substantially fewer computational resources to train



Results on ImageNet

	Ours-JFT (ViT-H/14)	Ours-JFT (ViT-L/16)	Ours-I21k (ViT-L/16)	BiT-L (ResNet152x4)	Noisy Student (EfficientNet-L2)
ImageNet	88.55 ± 0.04	87.76 ± 0.03	85.30 ± 0.02	87.54 ± 0.02	88.4/88.5*
ImageNet ReaL	90.72 ± 0.05	90.54 ± 0.03	88.62 ± 0.05	90.54	90.55
CIFAR-10	99.50 ± 0.06	99.42 ± 0.03	99.15 ± 0.03	99.37 ± 0.06	—
CIFAR-100	94.55 ± 0.04	93.90 ± 0.05	93.25 ± 0.05	93.51 ± 0.08	—
Oxford-IIIT Pets	97.56 ± 0.03	97.32 ± 0.11	94.67 ± 0.15	96.62 ± 0.23	—
Oxford Flowers-102	99.68 ± 0.02	99.74 ± 0.00	99.61 ± 0.02	99.63 ± 0.03	—
VTAB (19 tasks)	77.63 ± 0.23	76.28 ± 0.46	72.72 ± 0.21	76.29 ± 1.70	—
TPUv3-core-days	2.5k	0.68k	0.23k	9.9k	12.3k

Results on popular image classification benchmarks

Q & A

Web: www.personal.psu.edu/dux19/

Email: dux19@psu.edu